

Apache Sqoop Interview Questions and Answers

Q: What is Apache Sqoop?

A: A tool to transfer data between Hadoop and relational databases.

Q: What are the key features of Sqoop?

A: Parallel import/export, connectors for major RDBMS, incremental load, Hive/HBase integration, etc.

Q: What is the default file format used by Sqoop?

A: Delimited text file.

Q: What is the use of Sqoop import?

A: To import data from RDBMS to HDFS.

Q: What is the use of Sqoop export?

A: To export data from HDFS to RDBMS.

Q: What is incremental import in Sqoop?

A: Importing only new or updated data using `--incremental append` or `--incremental lastmodified`.

Q: What is the difference between append and lastmodified mode?

A: Append imports new rows; lastmodified imports updated rows.

Q: How do you perform a full import using Sqoop?

A: Without using the `--incremental` option.

Q: What is the use of `--split-by` in Sqoop?

A: Defines the column used to split data for parallel processing.

Q: What is the use of `--target-dir`?

A: Specifies the HDFS directory to store imported data.

Q: How can you import data into Hive using Sqoop?

A: Use `--hive-import` or `--hive-table`.

Q: How can you import data into HBase using Sqoop?

A: Use `--hbase-table` and `--column-family`.

Q: What is the use of `--direct` mode in Sqoop?

A: Enables database-specific fast import/export.