

# Untitled3

June 4, 2025

```
[4]: import pyspark
      from pyspark.sql import SparkSession
      from pyspark.sql.types import StructType, StructField, StringType, TimestampType
      from pyspark.sql.functions import col, to_timestamp
```

```
[5]: spark = SparkSession.builder\
      .master("local")\
      .appName("demo")\
      .getOrCreate()
```

```
[6]: schema = StructType([
      StructField("customer_nbr", StringType(), True),
      StructField("customer_desc", StringType(), True),
      StructField("start_ts", StringType(), True),
      StructField("end_ts", StringType(), True),
      StructField("create_ts", StringType(), True),
      StructField("last_update_ts", StringType(), True),
      StructField("client_id", StringType(), True)
    ])
```

```
[8]: # Load raw data from CSV with the defined schema
      df_raw = spark.read.csv("gs://newbucketrgcp1/customer_data_with_values.
      ↪csv", header = True, schema = schema)
      df_raw.show(0)
```

```
+-----+-----+-----+-----+-----+-----+-----+
|customer_nbr|customer_desc|start_ts|end_ts|create_ts|last_update_ts|client_id|
+-----+-----+-----+-----+-----+-----+-----+
only showing top 0 rows
```

```
[9]: df_raw.show()
```

[Stage 1:>

(0 + 1) / 1]

+-----+-----+-----+-----+-----+					
-----+-----+-----+-----+-----+					
customer_nbr customer_desc		start_ts		end_ts	
create_ts last_update_ts		client_id			
+-----+-----+-----+-----+-----+					
	CUST003	Customer C	2022-12-08 22:59:19	2023-11-28 13:23:03	2021-09-22
12:36:06		user3	2022-03-09 00:45:34		
	CUST004	Customer B	2022-09-23 17:31:40	2023-09-25 15:57:53	2021-03-08
11:52:34		user4	2022-08-31 17:57:57		
	CUST002	Customer D	2022-07-08 22:29:18	2023-02-24 03:16:23	2021-07-16
12:57:40		user2	2022-06-24 22:13:39		
	CUST004	Customer D	2022-12-25 20:41:29	2023-05-24 05:42:58	2021-05-30
19:28:21		user2	2022-06-10 18:07:29		
	CUST001	Customer D	2022-12-29 00:06:48	2023-06-03 16:54:36	2021-11-19
20:38:19		user4	2022-04-09 05:20:27		
	CUST004	Customer A	2022-01-20 01:03:28	2023-04-09 04:26:01	2021-12-01
22:47:42		user4	2022-07-02 22:05:32		
	CUST002	Customer C	2022-05-19 15:15:52	2023-09-04 21:57:20	2021-06-14
01:48:49		user4	2022-10-06 04:34:19		
	CUST004	Customer B	2022-09-18 05:45:43	2023-11-11 05:50:45	2021-11-20
11:08:28		user2	2022-12-12 11:56:58		
	CUST001	Customer B	2022-09-22 16:46:56	2023-12-24 09:50:40	2021-12-12
09:21:34		user4	2022-05-25 00:34:22		
	CUST003	Customer D	2022-11-29 23:45:45	2023-03-17 19:42:50	2021-10-20
17:05:13		user4	2022-10-30 07:01:22		
	CUST004	Customer D	2022-04-21 09:46:54	2023-02-15 05:14:24	2021-01-18
12:26:47		user1	2022-04-30 16:57:42		
	CUST004	Customer B	2022-02-26 17:59:11	2023-04-24 23:30:59	2021-02-04
18:24:10		user1	2022-04-01 03:32:01		
	CUST002	Customer C	2022-07-12 23:21:09	2023-12-27 01:53:30	2021-09-20
21:54:35		user3	2022-06-27 06:35:12		
	CUST001	Customer D	2022-10-02 13:36:42	2023-02-19 02:08:02	2021-12-06
00:09:33		user2	2022-04-15 04:45:35		
	CUST003	Customer D	2022-05-05 23:32:11	2023-08-25 11:08:05	2021-07-09
05:51:34		user2	2022-01-06 03:56:05		
	CUST003	Customer C	2022-04-25 09:21:10	2023-11-12 03:24:28	2021-08-01
06:01:45		user2	2022-05-30 17:08:02		
	CUST003	Customer C	2022-09-10 21:59:57	2023-02-07 05:01:21	2021-04-06
15:46:18		user3	2022-12-17 06:45:53		
	CUST001	Customer C	2022-01-05 05:16:43	2023-12-17 04:01:29	2021-12-26
07:47:14		user1	2022-12-06 17:04:53		
	CUST004	Customer C	2022-04-26 10:00:11	2023-05-04 21:14:12	2021-11-19
12:39:04		user1	2022-10-04 02:08:58		
	CUST001	Customer C	2022-04-29 21:19:24	2023-07-09 14:14:15	2021-09-23
17:39:17		user1	2022-07-06 04:23:33		
+-----+-----+-----+-----+-----+					
-----+-----+-----+-----+-----+					

only showing top 20 rows

```
[10]: # Transform timestamp columns from string to timestamp
df_transformed = df_raw \
    .withColumn("start_timestamp", to_timestamp(col("start_ts"), "yyyy-MM-dd HH:
    ↪mm:ss")) \
    .withColumn("end_timestamp", to_timestamp(col("end_ts"), "yyyy-MM-dd HH:mm:
    ↪ss")) \
    .withColumn("create_timestamp", to_timestamp(col("create_ts"), "yyyy-MM-dd
    ↪HH:mm:ss")) \
    .withColumn("last_update_timestamp", to_timestamp(col("last_update_ts"),
    ↪"yyyy-MM-dd HH:mm:ss")) \
    .drop("start_ts", "end_ts", "create_ts", "last_update_ts")
```

```
[11]: # Save transformed data to CSV
df_transformed.write.csv("transformed_data_struct1.csv", header=True)#,
    ↪mode="overwrite")

df_transformed.show(5)
```

```
+-----+-----+-----+-----+
+-----+-----+-----+-----+
|customer_nbr|customer_desc|      client_id|  start_timestamp|
end_timestamp|  create_timestamp|last_update_timestamp|
+-----+-----+-----+-----+
+-----+-----+-----+-----+
|      CUST003|  Customer C|2022-03-09 00:45:34|2022-12-08 22:59:19|2023-11-28
13:23:03|2021-09-22 12:36:06|          null|
|      CUST004|  Customer B|2022-08-31 17:57:57|2022-09-23 17:31:40|2023-09-25
15:57:53|2021-03-08 11:52:34|          null|
|      CUST002|  Customer D|2022-06-24 22:13:39|2022-07-08 22:29:18|2023-02-24
03:16:23|2021-07-16 12:57:40|          null|
|      CUST004|  Customer D|2022-06-10 18:07:29|2022-12-25 20:41:29|2023-05-24
05:42:58|2021-05-30 19:28:21|          null|
|      CUST001|  Customer D|2022-04-09 05:20:27|2022-12-29 00:06:48|2023-06-03
16:54:36|2021-11-19 20:38:19|          null|
+-----+-----+-----+-----+
+-----+-----+-----+-----+
only showing top 5 rows
```

```
[13]: df_raw1 = spark.read.csv("gs://newbucketrgcp1/employee_data.csv",header = True)
df_raw1.show(0)
```

```
+-----+-----+-----+-----+-----+
|empId|empName|empGender|empSalary|empCountry|
+-----+-----+-----+-----+-----+
only showing top 0 rows
```

```
[14]: df_raw1.show()
```

```
+-----+-----+-----+-----+-----+
|empId|empName|empGender|empSalary|empCountry|
+-----+-----+-----+-----+-----+
| 1|John Doe|Female|40256.214977607815|India|
| 2|John Doe|Female|54628.04698645289|UK|
| 3|Michael Johnson|Male|92119.45817408481|UK|
| 4|Emily Davis|Male|127470.53530973793|India|
| 5|John Doe|Male|131842.2807401374|Canada|
| 6|Jane Smith|Female|66457.25032866534|Australia|
| 7|Jane Smith|Male|97659.7503982054|UK|
| 8|John Doe|Male|102033.2220579916|Canada|
| 9|Jane Smith|Male|147863.5614979859|Canada|
| 10|Sam Brown|Male|135020.178183817|USA|
| 11|Sam Brown|Female|79012.57794462639|UK|
| 12|Michael Johnson|Male|49611.66080687332|India|
| 13|Emily Davis|Female|141959.93422365707|Canada|
| 14|Emily Davis|Male|115595.09477791714|Australia|
| 15|Michael Johnson|Male|39123.11478244402|UK|
| 16|Jane Smith|Male|58254.331529684576|USA|
| 17|Jane Smith|Male|120467.635616164|India|
| 18|Jane Smith|Male|75427.58315479774|Canada|
| 19|Jane Smith|Male|126288.83638980243|USA|
| 20|Emily Davis|Male|108338.68524293705|UK|
+-----+-----+-----+-----+-----+
only showing top 20 rows
```

```
[15]: from pyspark.sql.functions import col, explode, split, count

# Split the empName column into words and explode the result
words_df = df_raw1.select(explode(split(col("empName"), " ")).alias("word"))
```

```
# Group by word and count occurrences, then order by count in descending order
↳ and limit to top 3
top_words = words_df.groupBy("word").count().orderBy(col("count").desc()).
↳ limit(3)
top_words.show()
```

[Stage 9:>

(0 + 1) / 1]

```
+-----+-----+
| word|count|
+-----+-----+
| Jane|   10|
|Smith|   10|
|Davis|    9|
+-----+-----+
```

```
[16]: # Drop duplicate records from the DataFrame
df_no_duplicates = df_raw1.dropDuplicates()
df_no_duplicates.show()
```

[Stage 13:>

(0 + 1) / 1]

```
+-----+-----+-----+-----+-----+
| empId|      empName|empGender|      empSalary|empCountry|
+-----+-----+-----+-----+-----+
|   21|    Emily Davis|   Female| 87736.58716258111|    India|
|    8|      John Doe|    Male| 102033.2220579916|   Canada|
|   11|    Sam Brown|   Female| 79012.57794462639|     UK|
|   25|    Jane Smith|    Male| 120191.118217398| Australia|
|   15|Michael Johnson|    Male| 39123.11478244402|     UK|
|   20|    Emily Davis|    Male|108338.68524293705|     UK|
|    5|      John Doe|    Male| 131842.2807401374|   Canada|
|   26|    Emily Davis|   Female|126616.69717793733|   Canada|
|    2|      John Doe|   Female| 54628.04698645289|     UK|
|   39|Michael Johnson|   Female|40328.150028249336|   Canada|
|   29|    Sam Brown|    Male|113824.85958333808|   Canada|
|   17|    Jane Smith|    Male| 120467.635616164|    India|
|   13|    Emily Davis|   Female|141959.93422365707|   Canada|
|   38|    Sam Brown|   Female| 66819.14581953104|    India|
|   30|    Sam Brown|    Male| 42052.6967687742|     UK|
|   24|    Emily Davis|   Female| 63520.52748685201| Australia|
|    4|    Emily Davis|    Male|127470.53530973793|    India|
|   27|    Sam Brown|   Female|46661.455291065846|     USA|
|   37|      John Doe|    Male| 96903.92923049428|     UK|
|   28|    Sam Brown|   Female|104601.34691057618|     USA|
```

```
+-----+-----+-----+-----+-----+
only showing top 20 rows
```

```
[24]: # Split the empName column into words and explode the result
words_df = df_raw1.select(explode(split(col("empName"), " ")).alias("word"))

# Group by word and count occurrences
word_count = words_df.groupBy("word").count()
word_count.show()
```

[Stage 37:> (0 + 1) / 1]

```
+-----+-----+
|  word|count|
+-----+-----+
|  Davis|    9|
|  Smith|   10|
|Michael|    6|
|    Doe|    7|
|   John|    7|
|    Sam|    8|
|  Emily|    9|
|  Brown|    8|
|   Jane|   10|
|Johnson|    6|
+-----+-----+
```

```
[18]: from pyspark.sql.functions import avg

# Group by empCountry and calculate the average empSalary
avg_salary_by_country = df_raw1.groupBy("empCountry").agg(avg("empSalary").
    ↪alias("avg_salary"))
avg_salary_by_country.show()
```

[Stage 23:> (0 + 1) / 1]

```
+-----+-----+
|empCountry|    avg_salary|
+-----+-----+
|      India|75087.37664313955|
|       USA|91601.79658475956|
|       UK|77487.65026014007|
|   Canada|99339.65505965672|
|Australia|98809.73973473044|
```

```
+-----+-----+
```

```
[19]: # Drop rows with any null values
df_no_nulls = df_raw1.dropna()
df_no_nulls.show()

# Fill null values with a specific value
df_filled_nulls = df_raw1.fillna({"empSalary": 50000})
df_filled_nulls.show()
```

```
+-----+-----+-----+-----+-----+
|empId|      empName|empGender|      empSalary|empCountry|
+-----+-----+-----+-----+-----+
|  1|      John Doe|   Female|40256.214977607815|      India|
|  2|      John Doe|   Female| 54628.04698645289|        UK|
|  3|Michael Johnson|    Male| 92119.45817408481|        UK|
|  4|    Emily Davis|    Male|127470.53530973793|      India|
|  5|      John Doe|    Male| 131842.2807401374|    Canada|
|  6|    Jane Smith|   Female| 66457.25032866534|Australia|
|  7|    Jane Smith|    Male|  97659.7503982054|        UK|
|  8|      John Doe|    Male| 102033.2220579916|    Canada|
|  9|    Jane Smith|    Male| 147863.5614979859|    Canada|
| 10|    Sam Brown|    Male|  135020.178183817|       USA|
| 11|    Sam Brown|   Female| 79012.57794462639|        UK|
| 12|Michael Johnson|    Male| 49611.66080687332|      India|
| 13|    Emily Davis|   Female|141959.93422365707|    Canada|
| 14|    Emily Davis|    Male|115595.09477791714|Australia|
| 15|Michael Johnson|    Male| 39123.11478244402|        UK|
| 16|    Jane Smith|    Male|58254.331529684576|       USA|
| 17|    Jane Smith|    Male| 120467.635616164|      India|
| 18|    Jane Smith|    Male| 75427.58315479774|    Canada|
| 19|    Jane Smith|    Male|126288.83638980243|       USA|
| 20|    Emily Davis|    Male|108338.68524293705|        UK|
+-----+-----+-----+-----+-----+
```

only showing top 20 rows

```
+-----+-----+-----+-----+-----+
|empId|      empName|empGender|      empSalary|empCountry|
+-----+-----+-----+-----+-----+
|  1|      John Doe|   Female|40256.214977607815|      India|
|  2|      John Doe|   Female| 54628.04698645289|        UK|
|  3|Michael Johnson|    Male| 92119.45817408481|        UK|
|  4|    Emily Davis|    Male|127470.53530973793|      India|
|  5|      John Doe|    Male| 131842.2807401374|    Canada|
```

6	Jane Smith	Female	66457.25032866534	Australia
7	Jane Smith	Male	97659.7503982054	UK
8	John Doe	Male	102033.2220579916	Canada
9	Jane Smith	Male	147863.5614979859	Canada
10	Sam Brown	Male	135020.178183817	USA
11	Sam Brown	Female	79012.57794462639	UK
12	Michael Johnson	Male	49611.66080687332	India
13	Emily Davis	Female	141959.93422365707	Canada
14	Emily Davis	Male	115595.09477791714	Australia
15	Michael Johnson	Male	39123.11478244402	UK
16	Jane Smith	Male	58254.331529684576	USA
17	Jane Smith	Male	120467.635616164	India
18	Jane Smith	Male	75427.58315479774	Canada
19	Jane Smith	Male	126288.83638980243	USA
20	Emily Davis	Male	108338.68524293705	UK

+-----+-----+-----+-----+-----+

only showing top 20 rows

```
[20]: from pyspark.sql.functions import countDistinct
```

```
# Count distinct values in empCountry column
distinct_countries = df_raw1.select(countDistinct("empCountry").
    ↳alias("distinct_countries"))
distinct_countries.show()
```

[Stage 30:>

(0 + 1) / 1]

```
+-----+
|distinct_countries|
+-----+
|                    5|
+-----+
```

```
[21]: from pyspark.sql.functions import col
```

```
# Filter records where empSalary > 50,000
filtered_df = df_raw1.filter(col("empSalary") > 50000)
filtered_df.show()
```

```
+-----+-----+-----+-----+-----+
|empId|      empName|empGender|      empSalary|empCountry|
+-----+-----+-----+-----+-----+
```



2	John Doe	Female	54628.04698645289	UK
3	Michael Johnson	Male	92119.45817408481	UK
4	Emily Davis	Male	127470.53530973793	India
5	John Doe	Male	131842.2807401374	Canada
6	Jane Smith	Female	66457.25032866534	Australia
7	Jane Smith	Male	97659.7503982054	UK
8	John Doe	Male	102033.2220579916	Canada
9	Jane Smith	Male	147863.5614979859	Canada
10	Sam Brown	Male	135020.178183817	USA
11	Sam Brown	Female	79012.57794462639	UK
13	Emily Davis	Female	141959.93422365707	Canada
14	Emily Davis	Male	115595.09477791714	Australia
16	Jane Smith	Male	58254.331529684576	USA
17	Jane Smith	Male	120467.635616164	India
18	Jane Smith	Male	75427.58315479774	Canada
19	Jane Smith	Male	126288.83638980243	USA
20	Emily Davis	Male	108338.68524293705	UK
21	Emily Davis	Female	87736.58716258111	India
23	Sam Brown	Male	91526.81573411886	Australia
24	Emily Davis	Female	63520.52748685201	Australia

+-----+-----+-----+-----+-----+-----+

only showing top 20 rows

```
[26]: # Read a JSON file and convert it into a DataFrame
json_df = spark.read.json("gs://newbucketrgcp1/sample_json.json")
json_df.show()
```

_corrupt_record	duration	page
[  null  null		
{  null  null		
"user_id": "U...  null  null		
"name": "Alice",  null  null		
"location": {  null  null		
"city": "Ne...  null  null		
"country": ...  null  null		
},  null  null		
"sessions": [  null  null		
{  null  null		
"session_...  null  null		
"start_ti...  null  null		
"page_vie...  null  null		
null  30  /home		
null  120  /products		

```
|          ]|    null|    null|
|          },|    null|    null|
|          {|    null|    null|
|    "session_...|    null|    null|
|    "start_ti...|    null|    null|
```

```
+-----+-----+-----+
```

only showing top 20 rows

```
[27]: # Find the second highest salary
second_highest_salary = df_raw1.orderBy(col("empSalary").desc()).
    ↳select(col("empSalary")).distinct().collect()[1][0]
print(f"The second highest salary is: {second_highest_salary}")
```

[Stage 45:>

(0 + 1) / 1]

The second highest salary is: 91526.81573411886

```
[28]: # Sample data for the second DataFrame
data2 = [
    (1, "IT"),
    (2, "HR"),
    (3, "Finance"),
    (4, "Marketing"),
    (5, "Sales"),
    (6, "IT"),
    (7, "HR"),
    (8, "Finance"),
    (9, "Marketing"),
    (10, "Sales")
]

# Define the schema for the second DataFrame
schema2 = ["empId", "empDepartment"]

# Create the second DataFrame
df_raw2 = spark.createDataFrame(data2, schema2)

# Join two DataFrames and select specific columns
joined_df = df_raw1.join(df_raw2, df_raw1.empId == df_raw2.empId).
    ↳select(df_raw1.empId, df_raw1.empName, df_raw2.empDepartment)
joined_df.show()
```

[Stage 49:>

(0 + 1) / 1]

```
+-----+-----+-----+
```

empId	empName	empDepartment
1	John Doe	IT
2	John Doe	HR
3	Michael Johnson	Finance
4	Emily Davis	Marketing
5	John Doe	Sales
6	Jane Smith	IT
7	Jane Smith	HR
8	John Doe	Finance
9	Jane Smith	Marketing
10	Sam Brown	Sales

[ ]: