

Untitled5

June 5, 2025

```
[1]: from pyspark.sql import SparkSession
from pyspark.sql.functions import col, avg, countDistinct, desc, row_number,
    explode, split
from pyspark.sql.window import Window
```

```
[2]: spark = SparkSession.builder\
    .master("local")\
    .appName("demo")\
    .getOrCreate()
```

```
[3]: df = spark.read.option("header", True).csv("gs://datapyspark/employee_data.
    csv", inferSchema=True)
```

```
[4]: df.show()
```

| empId | empName | empGender | empSalary | empCountry | hireDate |
|-------|---------------------|-----------|-----------|----------------------|------------|
| 1 | Richard Williams | Female | 47329 | Malawi | 2018-05-05 |
| 2 | Devin Khan | Male | 103880 | Azerbaijan | 2015-12-23 |
| 3 | Mr. Michael Ayala | Female | 108197 | Rwanda | 2019-06-08 |
| 4 | Deborah Burnett | Male | 80102 | Colombia | 2022-10-10 |
| 5 | Heather Greene | Female | 96146 | Somalia | 2018-11-22 |
| 6 | Patricia Weiss | Female | 98525 | South Africa | 2021-06-21 |
| 7 | Nicholas Woodward | Male | 107971 | Jordan | 2020-11-15 |
| 8 | Amy Rodriguez | Female | 62957 | Vietnam | 2022-09-23 |
| 9 | Anthony Holt V | Male | 90547 | Marshall Islands | 2016-07-02 |
| 10 | Paul Wood | Male | 92935 | Bahrain | 2020-02-13 |
| 11 | Jeremy Johnson | Female | 77341 | United Arab Emirates | 2017-01-05 |
| 12 | John Brady | Female | 110007 | Slovenia | 2022-10-02 |
| 13 | Andrew Peterson | Female | 80136 | Colombia | 2016-08-08 |
| 14 | Dr. Bailey Anderson | Male | 75087 | Netherlands Antilles | 2016-06-30 |
| 15 | Paul Johnson | Male | 85158 | Belgium | 2021-03-08 |
| 16 | Daniel Rivera | Male | 87814 | Ukraine | 2020-11-07 |
| 17 | Evelyn Reed | Female | 115953 | Gibraltar | 2024-06-08 |
| 18 | David Newton | Female | 80115 | Uganda | 2023-12-24 |

| | | | | | |
|----|--------------|--------|--------|------------------|------------|
| 19 | Raymond Lee | Female | 87908 | Bermuda | 2024-08-20 |
| 20 | Kathryn Yang | Male | 108813 | Marshall Islands | 2023-03-30 |

only showing top 20 rows

```
[5]: # 1. Top 3 most occurring words in empName
words_df = df.select(explode(split(col("empName"), " ")).alias("word"))
top_words = words_df.groupBy("word").count().orderBy(desc("count")).limit(3)
top_words.show()
```

[Stage 5:=====> (1 + 1) / 3]

| word | count |
|----------|-------|
| Jennifer | 3 |
| Daniel | 3 |
| Paul | 3 |

```
[6]: # 2. Remove duplicate records
df_no_duplicates = df.dropDuplicates()
df_no_duplicates.show()
```

| empId | empName | empGender | empSalary | empCountry | hireDate |
|-------|------------------|-----------|-----------|----------------------|------------|
| 44 | Adrienne Sanchez | Female | 86290 | Liechtenstein | 2017-03-05 |
| 43 | Jaime Perez | Male | 67462 | New Caledonia | 2024-07-22 |
| 27 | Paul Nguyen | Female | 60067 | Netherlands | 2018-12-14 |
| 4 | Deborah Burnett | Male | 80102 | Colombia | 2022-10-10 |
| 25 | Denise Gonzales | Female | 58877 | Indonesia | 2018-01-10 |
| 2 | Devin Khan | Male | 103880 | Azerbaijan | 2015-12-23 |
| 28 | Allison Roberts | Female | 118978 | Sudan | 2018-09-25 |
| 39 | John Williams | Male | 61439 | Haiti | 2019-06-09 |
| 26 | Andrea Rose | Male | 93428 | Cyprus | 2016-10-05 |
| 33 | Amber Wright DVM | Female | 116155 | Bosnia and Herzeg... | 2019-10-20 |
| 10 | Paul Wood | Male | 92935 | Bahrain | 2020-02-13 |
| 31 | Ellen Baker | Female | 37040 | Spain | 2020-07-21 |
| 42 | Tiffany Deleon | Male | 48674 | Dominican Republic | 2018-09-19 |
| 29 | Jennifer Salinas | Male | 65388 | French Southern T... | 2019-07-18 |
| 9 | Anthony Holt V | Male | 90547 | Marshall Islands | 2016-07-02 |
| 21 | David Compton | Female | 71507 | Saint Kitts and N... | 2019-10-03 |
| 49 | Terry Mitchell | Female | 67579 | Guinea | 2019-02-06 |
| 32 | Kayla Nguyen | Male | 81046 | American Samoa | 2022-09-03 |

| | | | | | |
|----|--------------|--------|-------|-----------------|------------|
| 22 | Daniel Boone | Female | 95277 | Solomon Islands | 2016-05-13 |
| 36 | Jill Long | Male | 72665 | Bermuda | 2018-11-22 |

only showing top 20 rows

```
[7]: # 3. Word count using PySpark (on empName)
word_count = words_df.groupBy("word").count()
word_count.show()
```

```
[Stage 10:> (0 + 1) / 1]
```

| word | count |
|----------|-------|
| Compton | 1 |
| Andrea | 2 |
| Sanchez | 1 |
| Brady | 1 |
| James | 1 |
| Jaime | 1 |
| Gilmore | 1 |
| Denise | 1 |
| Nguyen | 2 |
| Jennifer | 3 |
| DVM | 1 |
| Rose | 1 |
| Williams | 2 |
| Rivera | 1 |
| Amber | 1 |
| Deborah | 1 |
| Peterson | 1 |
| Miller | 1 |
| Wright | 1 |
| Howell | 1 |

only showing top 20 rows

```
[8]: # 4. Group by empCountry and calculate average salary
avg_salary = df.groupBy("empCountry").agg(avg("empSalary").alias("avgSalary"))
avg_salary.show()
```

| empCountry | avgSalary |
|----------------------|-----------|
| Heard Island and ... | 39311.0 |

| | | |
|----------------------|--------------------|----------|
| French Southern T... | | 65388.0 |
| | Turkey | 66545.0 |
| | Malawi | 47329.0 |
| | Rwanda | 108197.0 |
| | Jordan | 107971.0 |
| | Sudan | 118978.0 |
| | Equatorial Guinea | 63383.0 |
| | Belgium | 85158.0 |
| | Ecuador | 62624.0 |
| | New Caledonia | 82432.0 |
| | American Samoa | 81046.0 |
| | Somalia | 96146.0 |
| | Marshall Islands | 99680.0 |
| Netherlands Antilles | | 75087.0 |
| | Spain | 37040.0 |
| | Russian Federation | 82395.0 |
| | Liechtenstein | 86290.0 |
| | Thailand | 46582.0 |
| | Ukraine | 87814.0 |

+-----+

only showing top 20 rows

```
[9]: # 5. Handle missing/null values
df_cleaned = df.na.fill({"empName": "Unknown", "empSalary": 0})
df_cleaned.show()
```

| empId | empName | empGender | empSalary | empCountry | hireDate |
|-------|---------|---------------------|-----------|------------|---------------------------------|
| | | | | | |
| | 1 | Richard Williams | Female | 47329 | Malawi 2018-05-05 |
| | 2 | Devin Khan | Male | 103880 | Azerbaijan 2015-12-23 |
| | 3 | Mr. Michael Ayala | Female | 108197 | Rwanda 2019-06-08 |
| | 4 | Deborah Burnett | Male | 80102 | Colombia 2022-10-10 |
| | 5 | Heather Greene | Female | 96146 | Somalia 2018-11-22 |
| | 6 | Patricia Weiss | Female | 98525 | South Africa 2021-06-21 |
| | 7 | Nicholas Woodward | Male | 107971 | Jordan 2020-11-15 |
| | 8 | Amy Rodriguez | Female | 62957 | Vietnam 2022-09-23 |
| | 9 | Anthony Holt V | Male | 90547 | Marshall Islands 2016-07-02 |
| | 10 | Paul Wood | Male | 92935 | Bahrain 2020-02-13 |
| | 11 | Jeremy Johnson | Female | 77341 | United Arab Emirates 2017-01-05 |
| | 12 | John Brady | Female | 110007 | Slovenia 2022-10-02 |
| | 13 | Andrew Peterson | Female | 80136 | Colombia 2016-08-08 |
| | 14 | Dr. Bailey Anderson | Male | 75087 | Netherlands Antilles 2016-06-30 |
| | 15 | Paul Johnson | Male | 85158 | Belgium 2021-03-08 |
| | 16 | Daniel Rivera | Male | 87814 | Ukraine 2020-11-07 |
| | 17 | Evelyn Reed | Female | 115953 | Gibraltar 2024-06-08 |
| | 18 | David Newton | Female | 80115 | Uganda 2023-12-24 |

| | | | | | |
|----|--------------|--------|--------|------------------|------------|
| 19 | Raymond Lee | Female | 87908 | Bermuda | 2024-08-20 |
| 20 | Kathryn Yang | Male | 108813 | Marshall Islands | 2023-03-30 |

only showing top 20 rows

```
[10]: # 6. Count distinct values in empCountry
distinct_countries = df.select("empCountry").distinct().count()
print(f"Distinct countries: {distinct_countries}")
```

Distinct countries: 45

```
[11]: # 7. Filter records where salary > 50000
high_salary = df.filter(col("empSalary") > 50000)
high_salary.show()
```

| empId | empName | empGender | empSalary | empCountry | hireDate |
|-------|---------------------|-----------|-----------|----------------------|------------|
| 2 | Devin Khan | Male | 103880 | Azerbaijan | 2015-12-23 |
| 3 | Mr. Michael Ayala | Female | 108197 | Rwanda | 2019-06-08 |
| 4 | Deborah Burnett | Male | 80102 | Colombia | 2022-10-10 |
| 5 | Heather Greene | Female | 96146 | Somalia | 2018-11-22 |
| 6 | Patricia Weiss | Female | 98525 | South Africa | 2021-06-21 |
| 7 | Nicholas Woodward | Male | 107971 | Jordan | 2020-11-15 |
| 8 | Amy Rodriguez | Female | 62957 | Vietnam | 2022-09-23 |
| 9 | Anthony Holt V | Male | 90547 | Marshall Islands | 2016-07-02 |
| 10 | Paul Wood | Male | 92935 | Bahrain | 2020-02-13 |
| 11 | Jeremy Johnson | Female | 77341 | United Arab Emirates | 2017-01-05 |
| 12 | John Brady | Female | 110007 | Slovenia | 2022-10-02 |
| 13 | Andrew Peterson | Female | 80136 | Colombia | 2016-08-08 |
| 14 | Dr. Bailey Anderson | Male | 75087 | Netherlands Antilles | 2016-06-30 |
| 15 | Paul Johnson | Male | 85158 | Belgium | 2021-03-08 |
| 16 | Daniel Rivera | Male | 87814 | Ukraine | 2020-11-07 |
| 17 | Evelyn Reed | Female | 115953 | Gibraltar | 2024-06-08 |
| 18 | David Newton | Female | 80115 | Uganda | 2023-12-24 |
| 19 | Raymond Lee | Female | 87908 | Bermuda | 2024-08-20 |
| 20 | Kathryn Yang | Male | 108813 | Marshall Islands | 2023-03-30 |
| 21 | David Compton | Female | 71507 | Saint Kitts and N... | 2019-10-03 |

only showing top 20 rows

```
[12]: # 9. Find second highest salary
window_spec = Window.orderBy(col("empSalary").desc())
second_highest_salary = df.withColumn("rank", row_number().over(window_spec)).
    filter(col("rank") == 2)
second_highest_salary.show()
```

```

25/06/05 10:57:04 WARN org.apache.spark.sql.execution.window.WindowExec: No
Partition Defined for Window operation! Moving all data to a single partition,
this can cause serious performance degradation.
25/06/05 10:57:04 WARN org.apache.spark.sql.execution.window.WindowExec: No
Partition Defined for Window operation! Moving all data to a single partition,
this can cause serious performance degradation.
25/06/05 10:57:05 WARN org.apache.spark.sql.execution.window.WindowExec: No
Partition Defined for Window operation! Moving all data to a single partition,
this can cause serious performance degradation.
25/06/05 10:57:05 WARN org.apache.spark.sql.execution.window.WindowExec: No
Partition Defined for Window operation! Moving all data to a single partition,
this can cause serious performance degradation.
[Stage 26:> (0 + 1) / 1]

```

```

+-----+-----+-----+-----+-----+-----+
+
|empId|      empName|empGender|empSalary|      empCountry|
hireDate|rank|
+-----+-----+-----+-----+-----+-----+
+
|   33|Amber Wright DVM|   Female|   116155|Bosnia and Herzeg...|2019-10-20|
2|
+-----+-----+-----+-----+-----+-----+
+

```

```

[13]: # 10. Join two DataFrames and select specific columns
df2 = df.select("empId", "empCountry").withColumnRenamed("empCountry", "country")
joined_df = df.join(df2, "empId").select("empId", "empName", "empSalary", "country")
joined_df.show()

```

```

+-----+-----+-----+-----+
|empId|      empName|empSalary|      country|
+-----+-----+-----+-----+
|   1|Richard Williams|   47329|      Malawi|
|   2|   Devin Khan|  103880|Azerbaijan|
|   3|Mr. Michael Ayala|  108197|      Rwanda|
|   4|Deborah Burnett|   80102|Colombia|
|   5|Heather Greene|   96146|      Somalia|
|   6|Patricia Weiss|   98525|South Africa|
|   7|Nicholas Woodward|  107971|      Jordan|
|   8|Amy Rodriguez|   62957|      Vietnam|
|   9|Anthony Holt V|   90547|Marshall Islands|
|  10|   Paul Wood|   92935|      Bahrain|
|  11|Jeremy Johnson|   77341|United Arab Emirates|

```

| | | | |
|----|---------------------|--------|----------------------|
| 12 | John Brady | 110007 | Slovenia |
| 13 | Andrew Peterson | 80136 | Colombia |
| 14 | Dr. Bailey Anderson | 75087 | Netherlands Antilles |
| 15 | Paul Johnson | 85158 | Belgium |
| 16 | Daniel Rivera | 87814 | Ukraine |
| 17 | Evelyn Reed | 115953 | Gibraltar |
| 18 | David Newton | 80115 | Uganda |
| 19 | Raymond Lee | 87908 | Bermuda |
| 20 | Kathryn Yang | 108813 | Marshall Islands |

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+

only showing top 20 rows

[16]: # 8. Read JSON file and convert to DataFrame

```
json_df = spark.read.option("multiline", True).json("gs://datapyspark/
↳sample_json.json")
json_df.show()
```

```
25/06/05 11:00:04 WARN org.apache.hadoop.util.concurrent.ExecutorHelper: Thread
(Thread[GetFileInfo #0,5,main]) interrupted:
java.lang.InterruptedException
    at
com.google.common.util.concurrent.AbstractFuture.get(AbstractFuture.java:510)
    at com.google.common.util.concurrent.FluentFuture$TrustedFuture.get(Flue
ntFuture.java:88)
    at org.apache.hadoop.util.concurrent.ExecutorHelper.logThrowableFromAfte
rExecute(ExecutorHelper.java:48)
    at org.apache.hadoop.util.concurrent.HadoopThreadPoolExecutor.afterExecu
te(HadoopThreadPoolExecutor.java:90)
    at
java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1157)
    at
java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:624)
    at java.lang.Thread.run(Thread.java:750)
25/06/05 11:00:04 WARN org.apache.hadoop.util.concurrent.ExecutorHelper: Thread
(Thread[GetFileInfo #1,5,main]) interrupted:
java.lang.InterruptedException
    at
com.google.common.util.concurrent.AbstractFuture.get(AbstractFuture.java:510)
    at com.google.common.util.concurrent.FluentFuture$TrustedFuture.get(Flue
ntFuture.java:88)
    at org.apache.hadoop.util.concurrent.ExecutorHelper.logThrowableFromAfte
rExecute(ExecutorHelper.java:48)
    at org.apache.hadoop.util.concurrent.HadoopThreadPoolExecutor.afterExecu
te(HadoopThreadPoolExecutor.java:90)
    at
java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1157)
    at
```

```
java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:624)
    at java.lang.Thread.run(Thread.java:750)
```

```
+-----+-----+-----+-----+
|      location|   name|      sessions|user_id|
+-----+-----+-----+-----+
| {New York, USA}| Alice|[ [{30, /home}, {...| U001|
| {San Francisco, USA}| Bob|[ [{10, /home}, {...| U002|
| {Madrid, Spain}| Carlos|      []| U003|
| {Berlin, Germany}| Diana|[ [{20, /home}, {...| U004|
| {Toronto, Canada}| Ethan|[ [{10, /login}],...| U005|
| {Dubai, UAE}| Fatima|      []| U006|
| {London, UK}| George|[ [{60, /home}], ...| U007|
| {Tokyo, Japan}| Hiroshi|[ [{15, /home}, {...| U008|
| {Paris, France}| Isabelle|      []| U009|
| {Melbourne, Austr...| Jack|[ [{5, /login}, {...| U010|
+-----+-----+-----+-----+
```

[]: